

These instructions are also available at:



Accessing NERSC's Systems

You have two options: Jupyter or SSH:

1. (jupyter) Go to: <https://jupyter.nersc.gov/> and sign in with your credentials from step 2
2. (ssh) In a Unix/Linux terminal, type:
elvis@laptop> ssh < user>@perlmutter.nersc.gov



Accessing the Terminal from within Jupyter

You might not want to install a terminal emulator and ssh – in that case, you can access a terminal window on jupyter.nersc.gov (step 3 above) as follows:

3. After logging into jupyter.nersc.gov (step 3) you should see something like this:

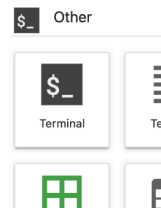
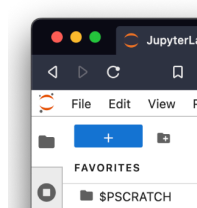
	Login Node	Shared GPU Node	Exclusive CPU Node	Exclusive GPU Node	Configurable Job
Alvarez	<button>start</button>	<button>start</button>	<button>start</button>	<button>start</button>	
Muller	<button>start</button>	<button>start</button>	<button>start</button>	<button>start</button>	<button>start</button>
Perlmutter	<button>stop</button> <button>server</button>	<button>start</button>	<button>start</button>	<button>start</button>	<button>start</button>
Resources	Use a login node shared with other users, outside the batch queues.	Use a single GPU on a node within a job allocation using defaults.	Use your own node within a job allocation using defaults.	Use multiple compute nodes with specialized settings.	
Use Cases	Visualization and analytics that are not memory intensive and can run on just a few cores.	Work that fits on a single GPU, and uses at most a quarter of a GPU node's CPU cores and host memory.	Visualization, analytics, machine learning that is compute or memory intensive but can be done on a single node.	Multi-node analytics jobs, jobs in reservations, custom project charging, and more.	

(You might not see the bright red “stop” button, and probably fewer rows/columns – that’s OK)

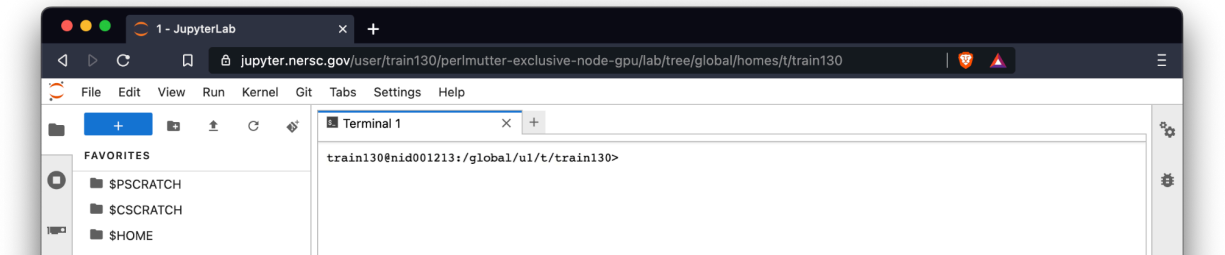
4. Select “Server” in the “Login Node” column and “Perlmutter” row

Login Node	
Alvarez	<button>start</button>
Muller	<button>start</button>
Perlmutter	<button>stop</button> <button>server</button>
Resources	Use a login node shared with other users, outside the batch queues.

5. After a short while, you should see a blue button (with a “+” sign) in the top left hand corner. Push it, and then select “Terminal” (you might need to scroll)



6. If you did everything correctly, you should see a terminal window in the left-hand tab:



Installing the Jupyter Kernels at NERSC

This tutorial requires specialized kernels to be installed – that's an automated process which you need to initiate in your training account

7. Access a shell on Perlmutter (either step 4, or step 5-8)
8. Clone the tutorial repository:

```
elvis@login12> git clone  
git@github.com:JuliaParallel/julia-hpc-tutorial-lan125.git
```
9. Enter the tutorial folder:

```
elvis@login12> cd julia-hpc-tutorial-lan125
```
10. Run the install script:

```
elvis@login12> ./install.sh
```

(this might take some time – that's OK)

Running Gray-Scott on Perlmutter/NERSC

These instructions are for the [Perlmutter system at NERSC](#) and are based on the [ORNL Gray Scott Tutorial](#) (qr code after step 18, below)



Configuring Gray-Scott on Perlmutter

13. Access a shell on Perlmutter (either step 4, or step 5-8)
14. Obtain Gray-Scott from GitHub, first access your scratch area and create a user-specific directory, clone the repository pointing at the GrayScott-JACC branch.

```
14.1 cd $SCRATCH
14.2 mkdir $USER
14.3 cd $USER
14.4 git clone --branch GrayScott-JACC
https://github.com/JuliaORNL/GrayScott.jl.git
```

15. Run the script prepared for this tutorial

[GrayScott.jl/scripts/config_perlmutter.sh](#)

to set up modules, environment, and packages.

```
source GrayScott.jl/scripts/config_perlmutter.sh
```



Running Gray-Scott jobs on Perlmutter

16. Create an area for Gray-Scott runs outside the repository (e.g. run001, future runs will be in run002, run003, etc.)

```
mkdir run001
```

17. Copy the Gray-Scott settings file and the job_perlmutter.sh to the run directory

```
17.1 cp GrayScott.jl/examples/settings-files.json run001
```

```
17.2 cp GrayScott.jl/scripts/job_perlmutter.sh run001
```

18. Submit your first job to Perlmutter. It should generate an adios bp file output, and total runtime should be around 12 seconds using a single MPI process and NVIDIA GPU.

```
18.1 cd run001
```

```
18.2 sbatch job_perlmutter.sh
```

To validate the output, please refer to the [ORNL Gray Scott Tutorial](#)



Slurm Jobscript Cheatsheet

These are the parameters you need to give to Slurm in order to submit jobs as a training account user.

Note: the GPU node reservation goes from 1:30pm - 5pm ET, Sun Nov 17th. During the reservation you should use:

19. To use 1 GPU only (sample flags for sbatch or salloc) use Slurm parameters

```
-A ntrain1 --reservation=lanl_training -C gpu -N 1 -c 32 -G 1 -t
30:00 -q shared
```

20. To use multiple nodes (sample flags for sbatch or salloc) use Slurm parameters:

```
-A ntrain1 --reservation=lanl_training -C gpu -N 2 -t 30:00 -q  
regular
```

Outside of the reservation, you should use:

21. To use 1 GPU only (sample flags for sbatch or salloc) use Slurm parameters

```
-A ntrain1 -C gpu -N 1 -c 32 -G 1 -t 30:00 -q shared
```

22. To use multiple nodes (sample flags for sbatch or salloc) use Slurm parameters:

```
-A ntrain1 -C gpu -N 2 -t 30:00 -q regular
```